


pyspark-pictures data frames



Learn the pyspark API through pictures and simple examples

<https://github.com/jkthompson/pyspark-pictures>

data frame



row



user input


groupby
function

aggregate
function

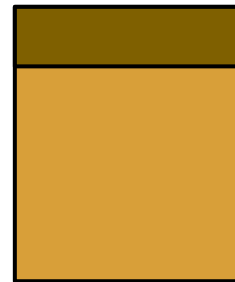
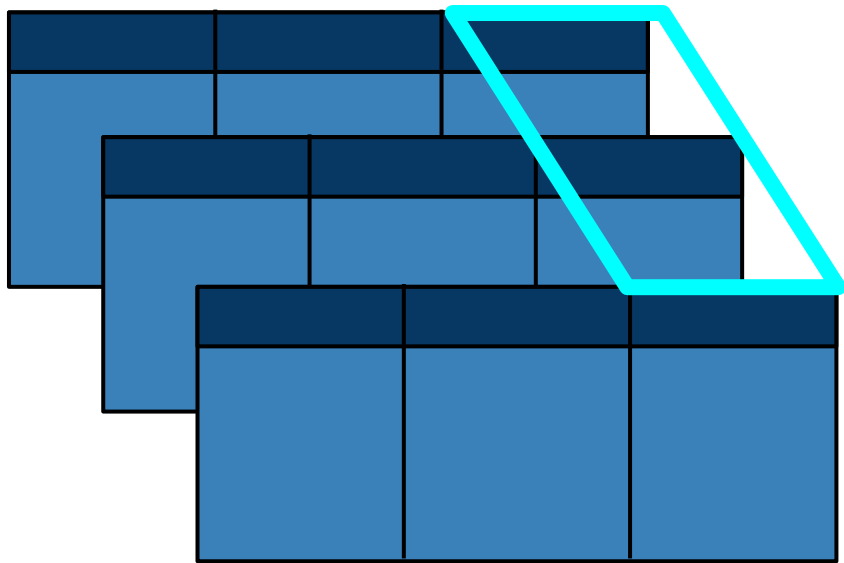
user function

spark input

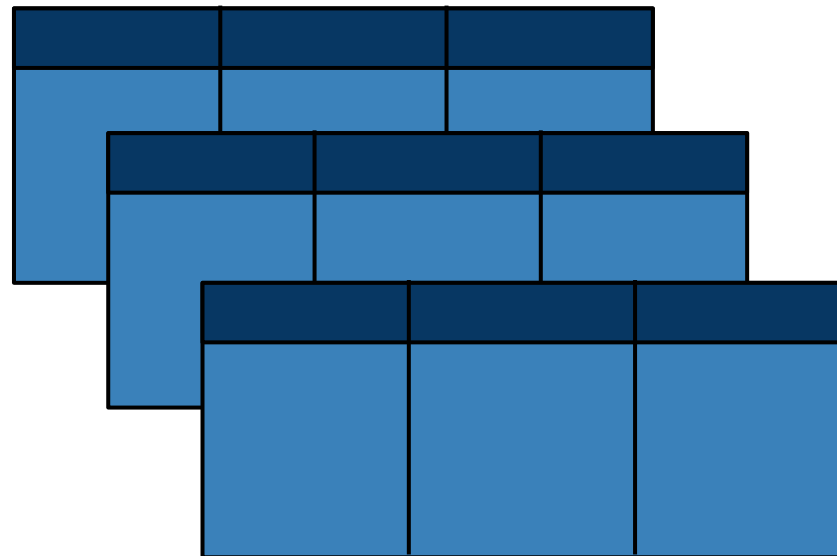
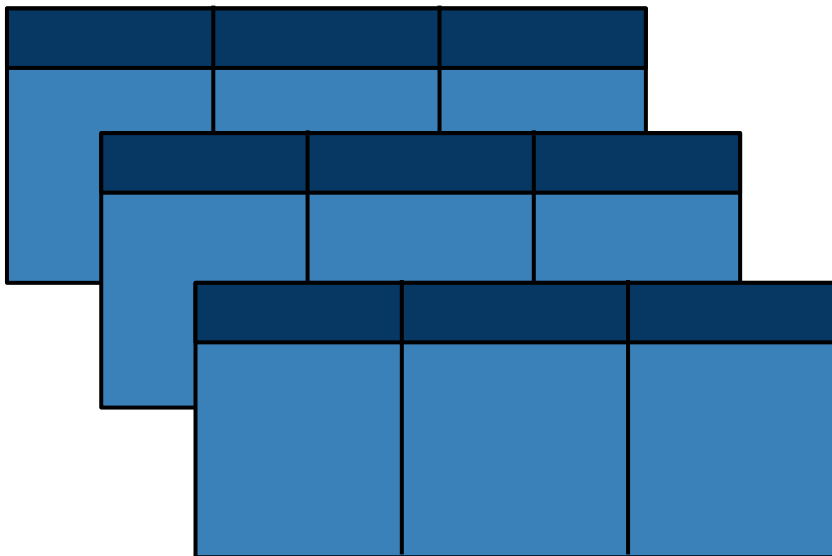
output



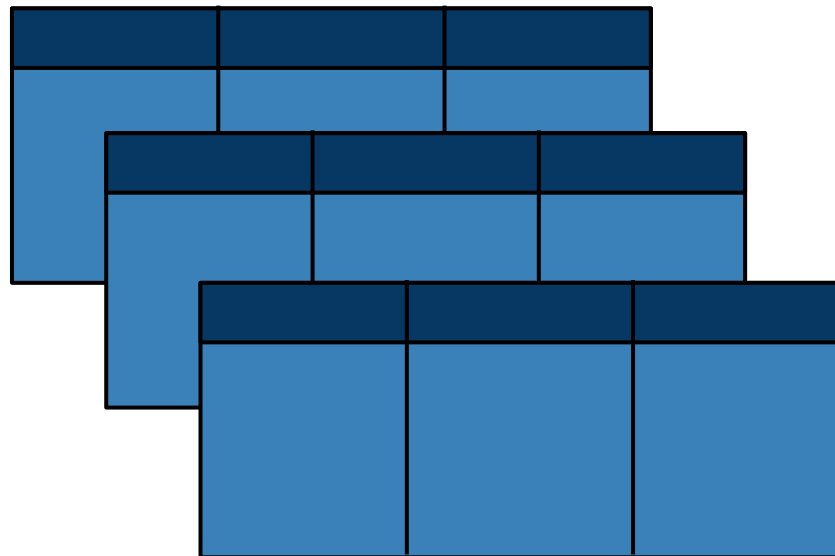
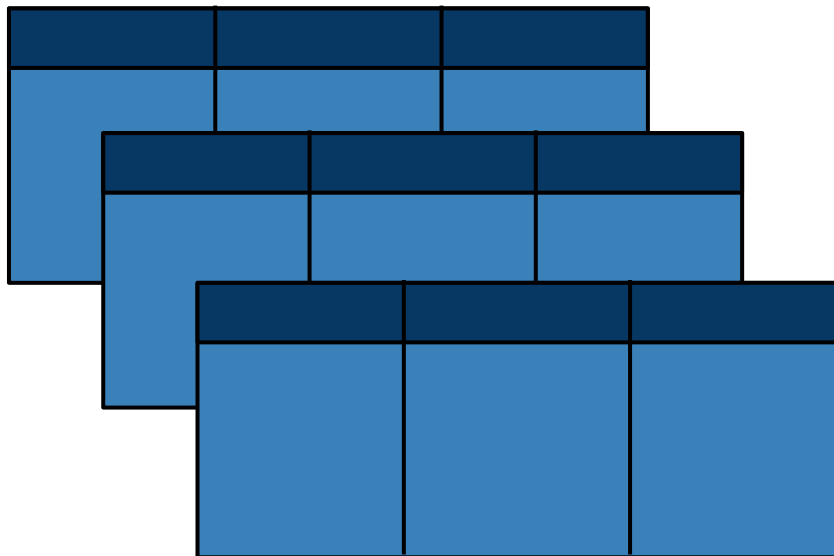
agg



alias

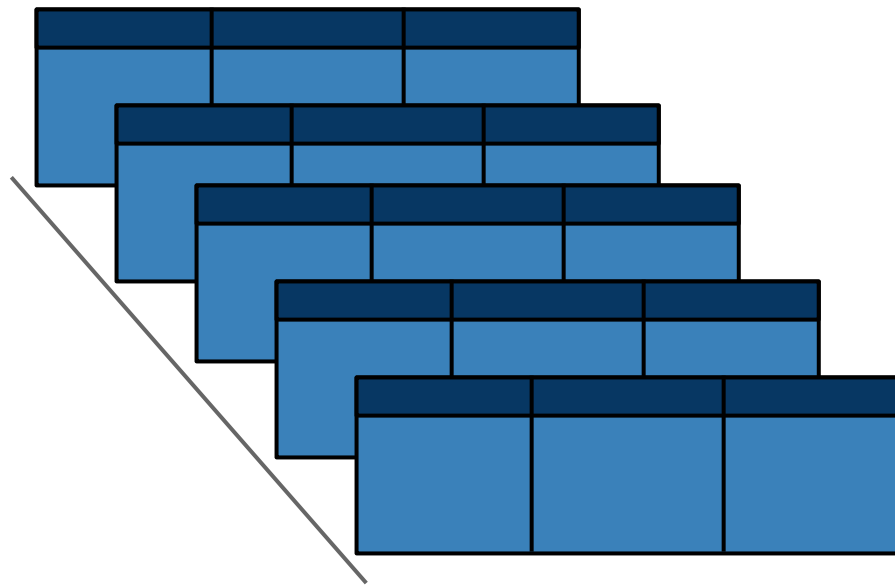
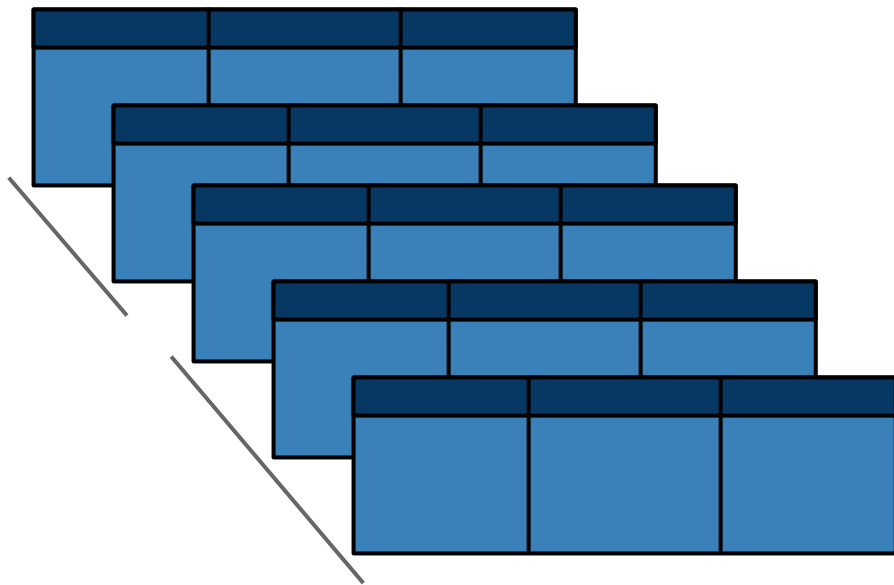


cache

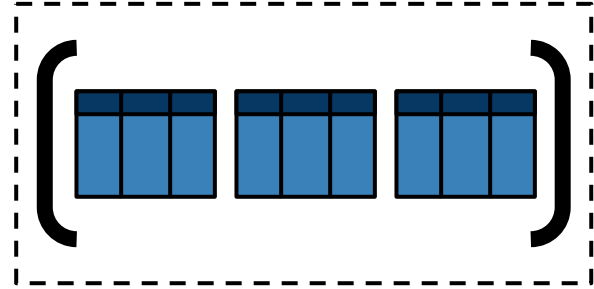



coalesce

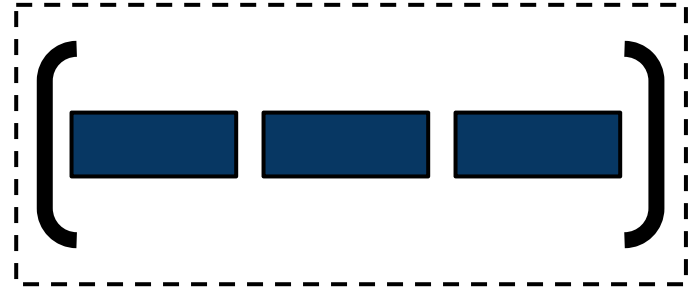
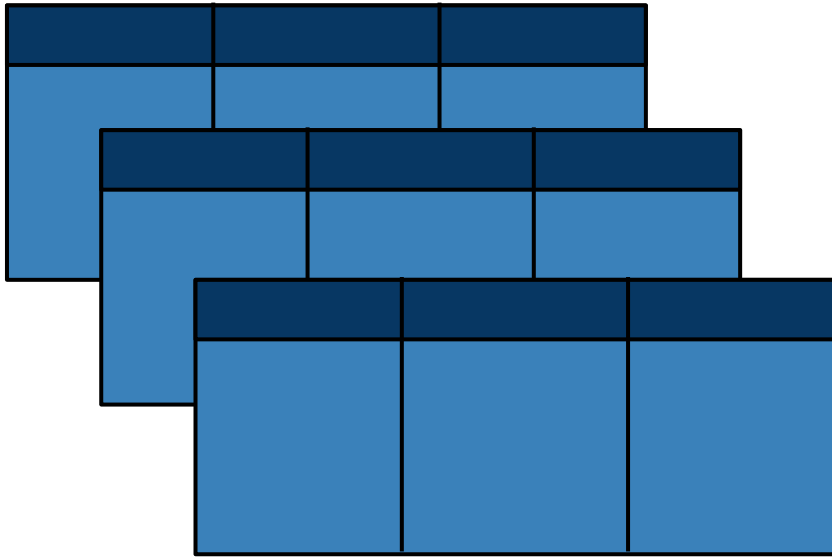
numPartitions = 1




collect



columns



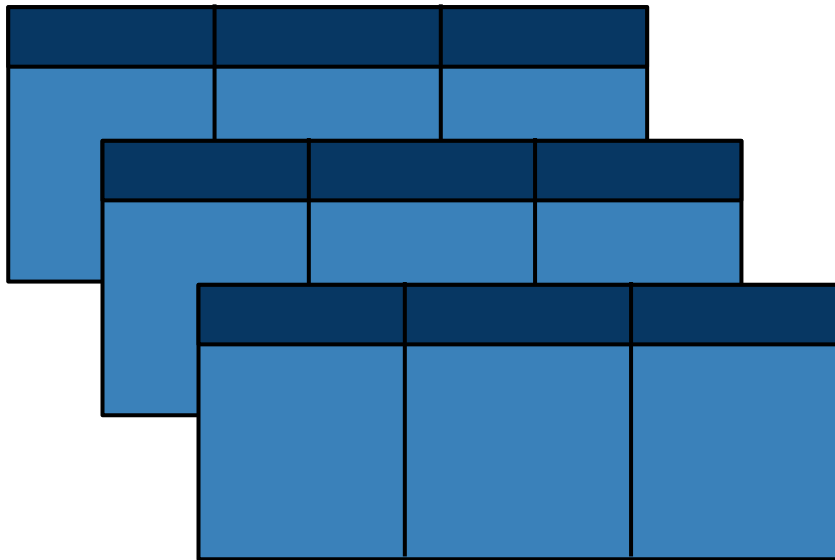
corr



Pearson's r

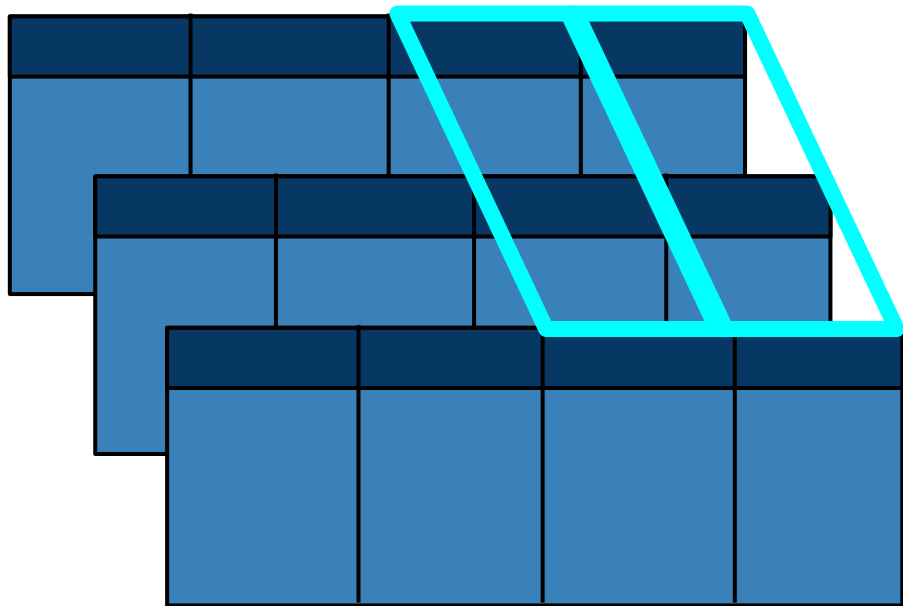
$$r = \frac{\sum_i (A_i - \bar{A})(C_i - \bar{C})}{\sqrt{\sum_i (A_i - \bar{A})^2} \sqrt{\sum_i (C_i - \bar{C})^2}}$$

count



3

COV

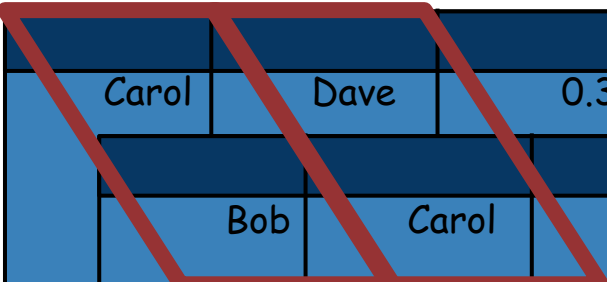


Sample Covariance

$$\frac{1}{N-1} \sum_i (A_i - \bar{A})(C_i - \bar{C})$$

crosstab

col1 = 'from' col2 = 'to'



| Carol | Dave | 0.3 |
|-------|-------|-----|
| Bob | Carol | 0.2 |
| from | to | amt |
| Alice | Bob | 0.1 |

| Carol | 0 | 0 | 1 |
|---------|-----|-------|------|
| Bob | 0 | 1 | 0 |
| from_to | Bob | Carol | Dave |
| Alice | 1 | 0 | 0 |

cube

*cols = 'from', 'to'

| Alice | Carol | 0.2 |
|-------|-------|-----|
| from | to | amt |
| Alice | Bob | 0.1 |

| null | null | |
|-------|-------|----------|
| Alice | null | |
| null | Carol | |
| null | Bob | |
| Alice | Carol | |
| from | to | agg(amt) |
| Alice | Bob | |

describe

| Carol | Dave | 0.3 |
|-------|-------|-----|
| Bob | Carol | 0.2 |
| from | to | amt |
| Alice | Bob | 0.1 |

| max | |
|---------|-----|
| | |
| min | |
| | |
| stdev | |
| | |
| mean | |
| summary | amt |
| count | |

distinct

| Bob | Carol | 0.2 |
|-------|-------|-----|
| | | |
| Carol | Dave | 0.3 |
| | | |
| Bob | Carol | 0.2 |
| | | |
| from | to | amt |
| Alice | Bob | 0.1 |

| Carol | Dave | 0.3 |
|-------|-------|-----|
| | | |
| Bob | Carol | 0.2 |
| | | |
| from | to | amt |
| Alice | Bob | 0.1 |

drop

col = 'amt'

| Carol | Dave | 0.3 |
|-------|-------|-----|
| | | |
| Bob | Carol | 0.2 |
| from | to | amt |
| Alice | Bob | 0.1 |

| Carol | Dave |
|-------|-------|
| | |
| Bob | Carol |
| from | to |
| Alice | Bob |

dropDuplicates

subset = ['from', 'to']

| Bob | Carol | 0.2 |
|-------|-------|-----|
| | | |
| Bob | Carol | 0.3 |
| | | |
| Bob | Carol | 0.2 |
| from | to | amt |
| Alice | Bob | 0.1 |

| Bob | Carol | 0.2 |
|-------|-------|-----|
| from | to | amt |
| Alice | Bob | 0.1 |

dropna

how = 'any' subset = ['from', 'to']

| Bob | Carol | 0.2 |
|-------|-------|------|
| | | |
| Carol | null | 0.3 |
| | | |
| Bob | Carol | null |
| from | to | amt |
| null | Bob | 0.1 |

| Bob | Carol | 0.2 |
|------|-------|------|
| from | to | amt |
| Bob | Carol | null |

dtypes

| Carol | Dave | 0.3 |
|-------|-------|-----|
| | | |
| Bob | Carol | 0.2 |
| from | to | amt |
| Alice | Bob | 0.1 |

[('from','string'), ('to','string'), ('amt','double')]

explain

extended = True

| Carol | Dave | 0.3 |
|-------|------|-----|

| Bob | Carol | 0.2 |
|-----|-------|-----|

| from | to | amt |
|-------|-----|-----|
| Alice | Bob | 0.1 |

== Parsed Logical Plan ==

...

== Analyzed Logical Plan ==

...

== Optimized Logical Plan ==

...

== Physical Plan ==

...

== RDD ==

fillna

value = 'unknown' subset = ['from', 'to']

| Carol | null | 0.3 |
|-------|------|-----|

| Bob | Carol | nan |
|-----|-------|-----|

| from | to | amt |
|------|-----|-----|
| null | Bob | 0.1 |

| Carol | unknown | 0.3 |
|-------|---------|-----|

| Bob | Carol | nan |
|-----|-------|-----|

| from | to | amt |
|---------|-----|-----|
| unknown | Bob | 0.1 |

filter

condition = "amt > 0.1"

| Carol | Dave | 0.3 |
|-------|-------|-----|
| | | |
| Bob | Carol | 0.2 |
| from | to | amt |
| Alice | Bob | 0.1 |

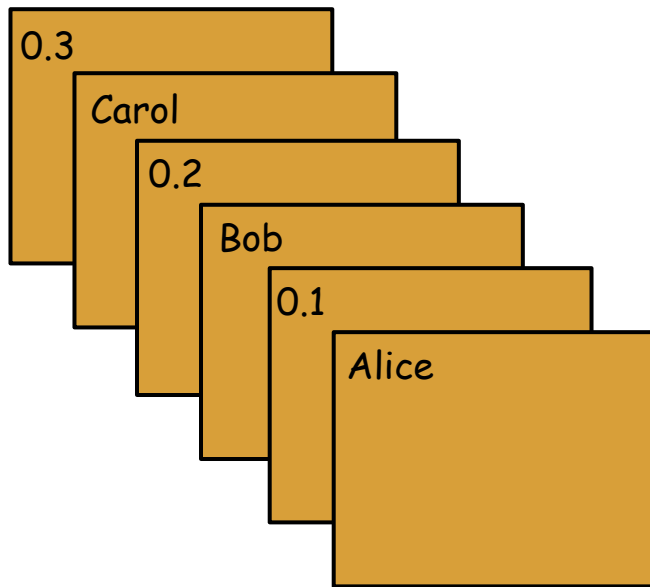

| Carol | Dave | 0.3 |
|-------|-------|-----|
| from | to | amt |
| Bob | Carol | 0.2 |

first

| Carol | Dave | 0.3 |
|-------|-------|-----|
| | | |
| Bob | Carol | 0.2 |
| from | to | amt |
| Alice | Bob | 0.1 |

Row(from='Alice', to='Bob', amt=0.1)

flatMap



foreach

| | | |
|-------|------|-----|
| | | |
| Carol | Dave | 0.3 |

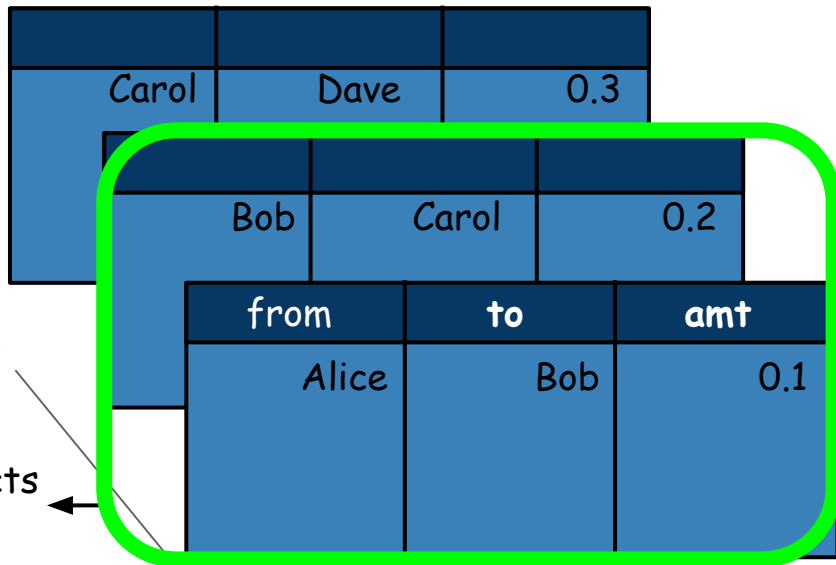
| | | |
|-----|-------|-----|
| | | |
| Bob | Carol | 0.2 |

| from | to | amt |
|-------|-----|-----|
| Alice | Bob | 0.1 |

side effects
(e.g print) ←

*no return value,
original DataFrame
unchanged

foreachPartition



The diagram illustrates the `foreachPartition` operation. It shows three nested DataFrames. The outermost DataFrame has a header row and one data row (Carol, Dave, 0.3). The middle DataFrame also has a header row and one data row (Bob, Carol, 0.2). The innermost DataFrame has a header row and one data row (Alice, Bob, 0.1). A green rounded rectangle highlights the innermost DataFrame, indicating it is the current partition being processed. An arrow points from the text 'side effects (e.g print)' to the green box.

| Carol | Dave | 0.3 |
|-------|------|-----|

| Bob | Carol | 0.2 |
|-----|-------|-----|

| from | to | amt |
|-------|-----|-----|
| Alice | Bob | 0.1 |

side effects
(e.g print)

*no return value,
original DataFrame
unchanged


freqItems

cols = ['from', 'amt'] support = 0.8

| Carol | Bob | 0.1 |
|-------|-------|-----|
| Alice | Bob | 0.5 |
| Alice | Bob | 0.1 |
| Alice | Dave | 0.1 |
| from | to | amt |
| Bob | Carol | 0.1 |

| from_freqItems | amt_freqItems |
|----------------|---------------|
| [Alice] | [0.1] |


groupBy (groupby)



GroupedData Object
with methods: *agg, avg, count,*
max, mean, min, pivot, sum

groupBy(col1).avg(col2)

col1 = 'from' col2 = 'amt'



| Carol | Dave | 0.3 |
|-------|-------|-----|
| Alice | Carol | 0.2 |
| from | to | amt |
| Alice | Bob | 0.1 |

| Carol | 0.3 |
|-------|----------|
| from | avg(amt) |
| Alice | 0.15 |

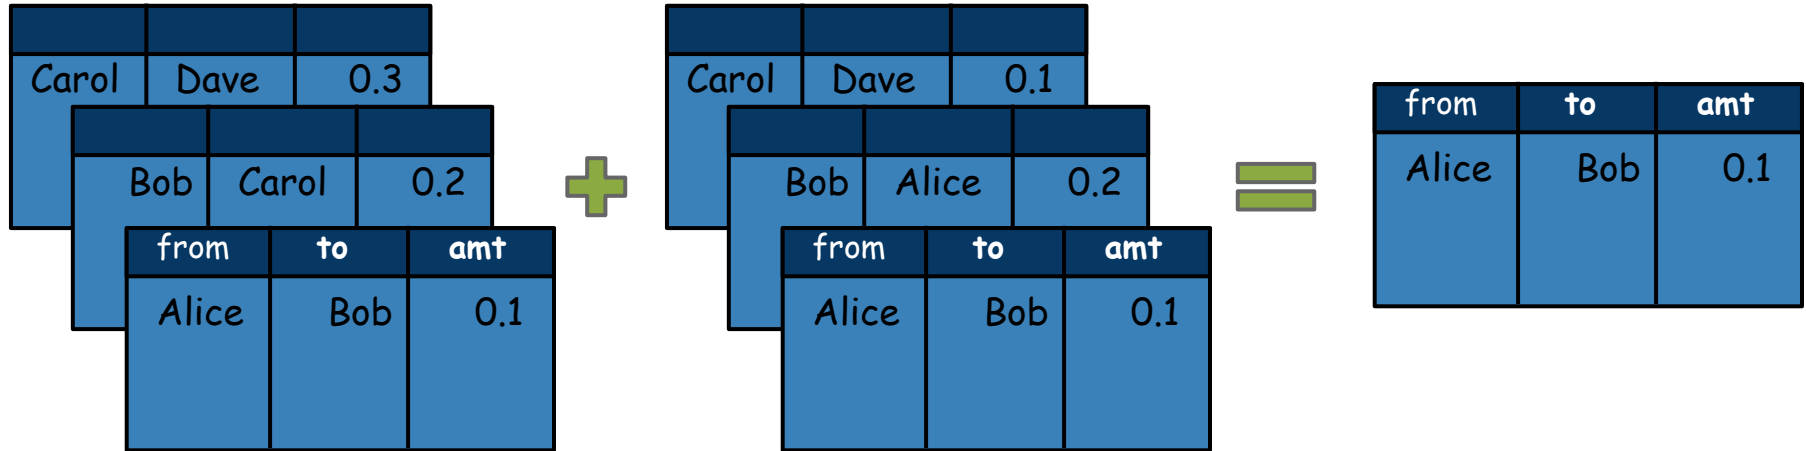
head

n = 2

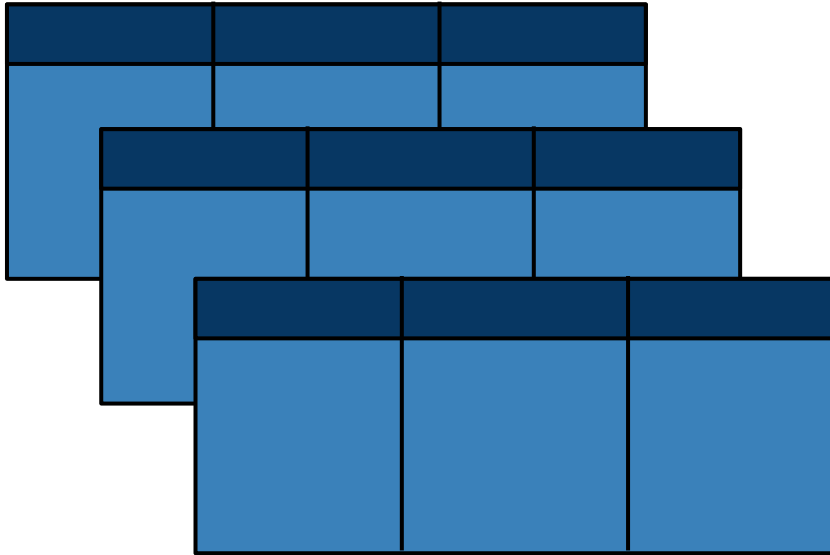
| Carol | Dave | 0.3 |
|-------|-------|-----|
| | | |
| Bob | Carol | 0.2 |
| from | to | amt |
| Alice | Bob | 0.1 |

[Row(from=u'Alice', to=u'Bob', amt=0.1),
Row(from=u'Bob', to=u'Carol', amt=0.2)]

intersect



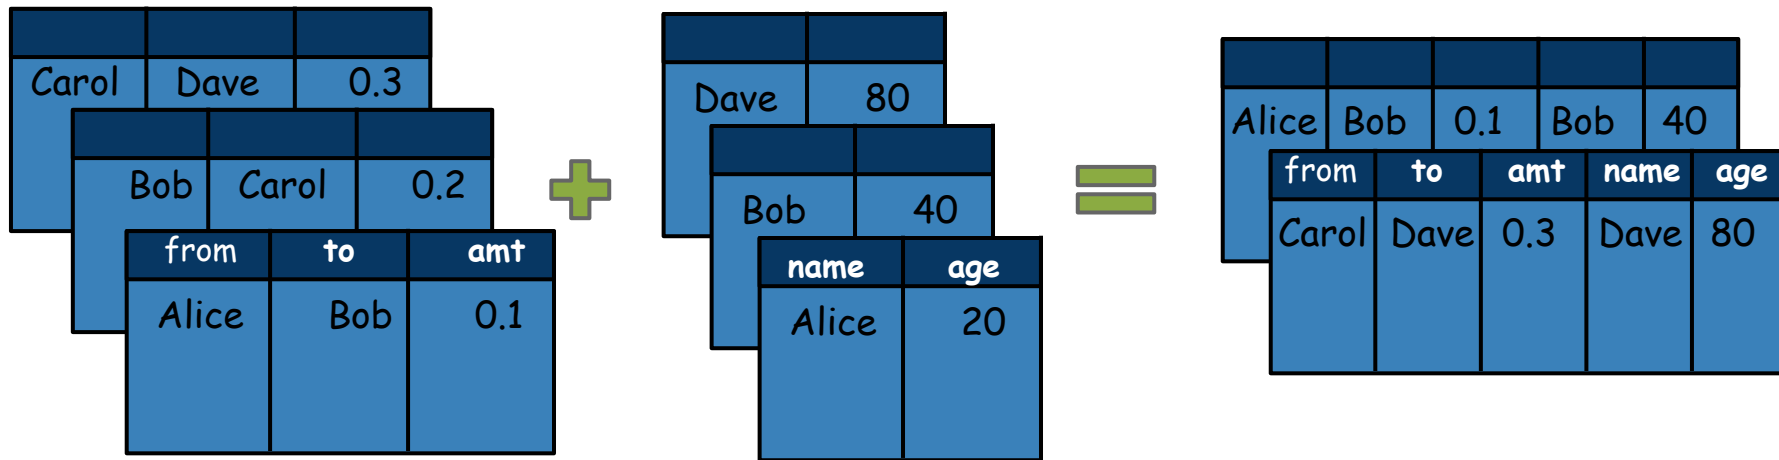
isLocal



False



join

joinExprs = x.to==y.name joinType = 'inner'

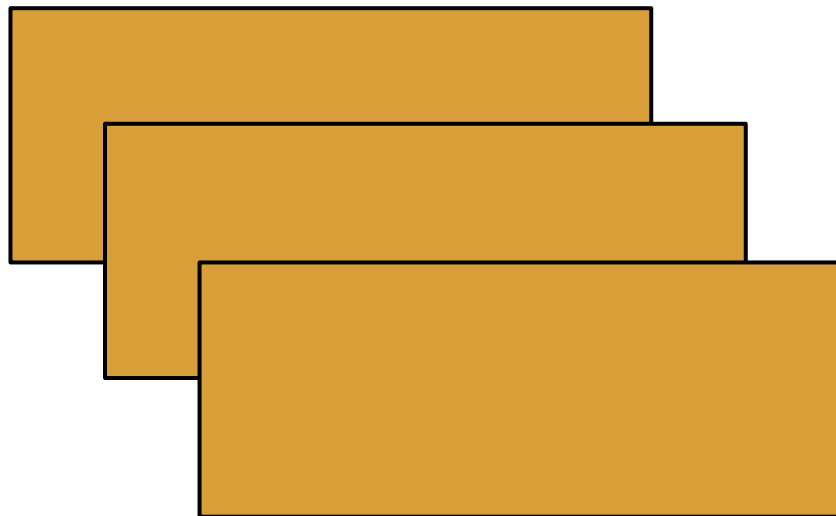


limit


num = 2



map



mapPartitions



na

| Bob | Carol | 0.2 |
|-----|-------|-----|

| Carol | null | 0.3 |
|-------|------|-----|

| Bob | Carol | null |
|-----|-------|------|

| from | to | amt |
|------|-----|-----|
| null | Bob | 0.1 |

DataFrameNaFunctions
Object

with methods: drop, fill, replace

orderBy


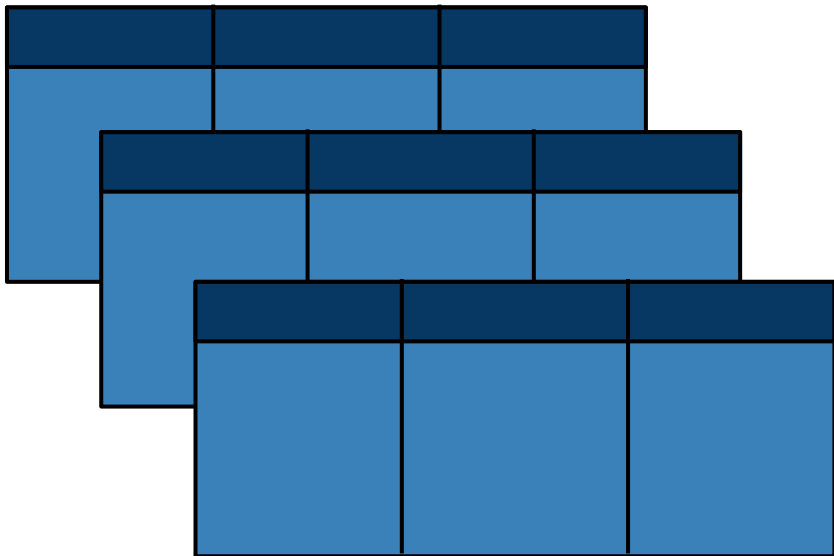
```
cols = ['from'], ascending = [False]
```

| | | |
|-------|-------|-----|
| | | |
| Carol | Dave | 0.3 |
| | | |
| Bob | Carol | 0.2 |
| | | |
| from | to | amt |
| Alice | Bob | 0.1 |

| | | | |
|-------|-------|-------|-----|
| | | | |
| Alice | Bob | 0.1 | |
| | | | |
| | Bob | Carol | 0.2 |
| | | | |
| | from | to | amt |
| | Carol | Dave | 0.3 |
| | | | |

persist

```
storageLevel =  
StorageLevel(MEMORY_ONLY_SER)
```



printSchema

| | | |
|-------|------|-----|
| | | |
| Carol | Dave | 0.3 |

| | | |
|-----|-------|-----|
| | | |
| Bob | Carol | 0.2 |

| from | to | amt |
|-------|-----|-----|
| Alice | Bob | 0.1 |

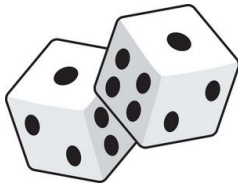
stdout

root

|-- from: string (nullable = true)
|-- to: string (nullable = true)
|-- amt: double (nullable = true)

randomSplit

weights = [0.5,0.5]

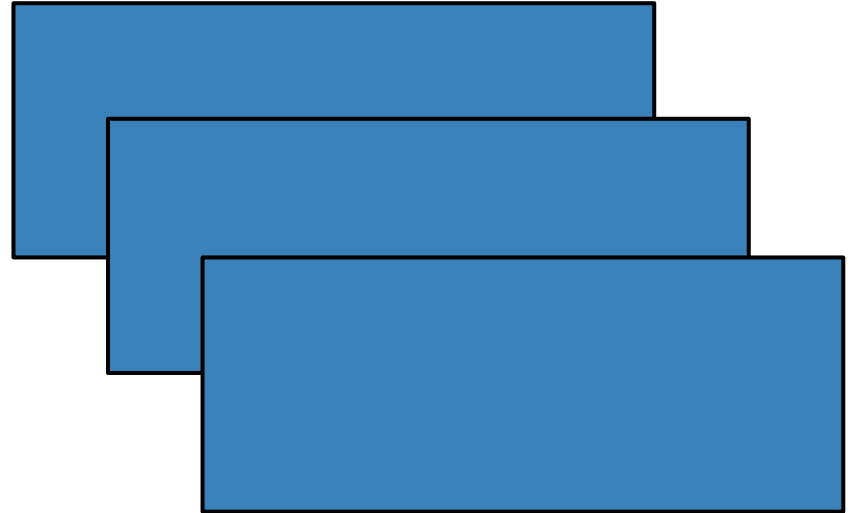
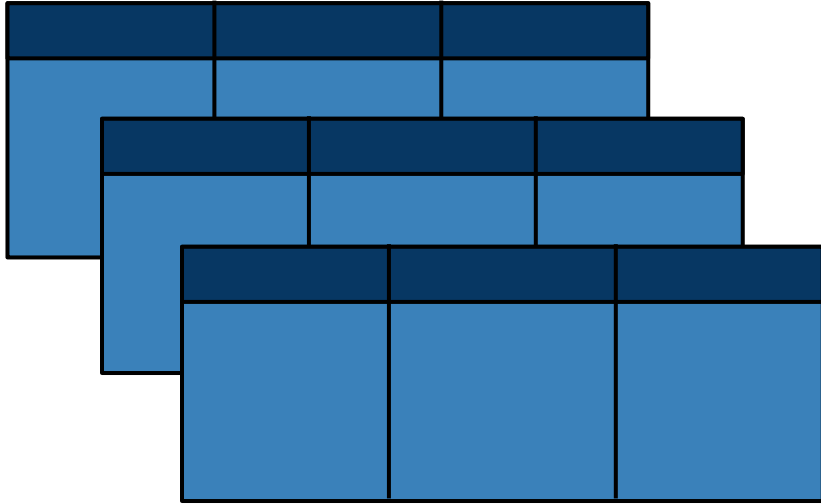


| Carol | Dave | 0.3 |
|-------|-------|-----|
| | | |
| Bob | Carol | 0.2 |
| from | to | amt |
| Alice | Bob | 0.1 |

| Carol | Dave | 0.3 |
|-------|-------|-----|
| from | to | amt |
| Bob | Carol | 0.2 |

| from | to | amt |
|-------|-----|-----|
| Alice | Bob | 0.1 |

rdd



registerTempTable

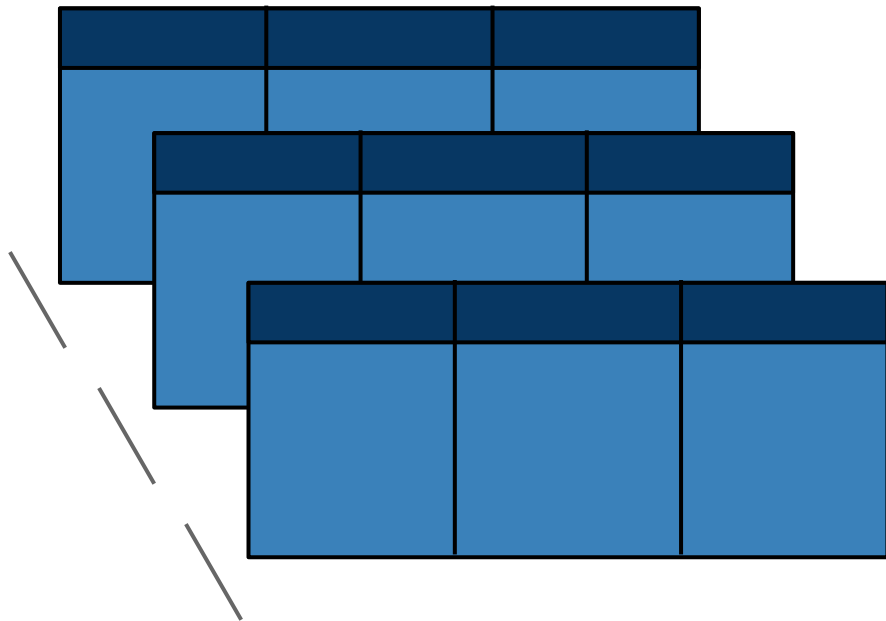
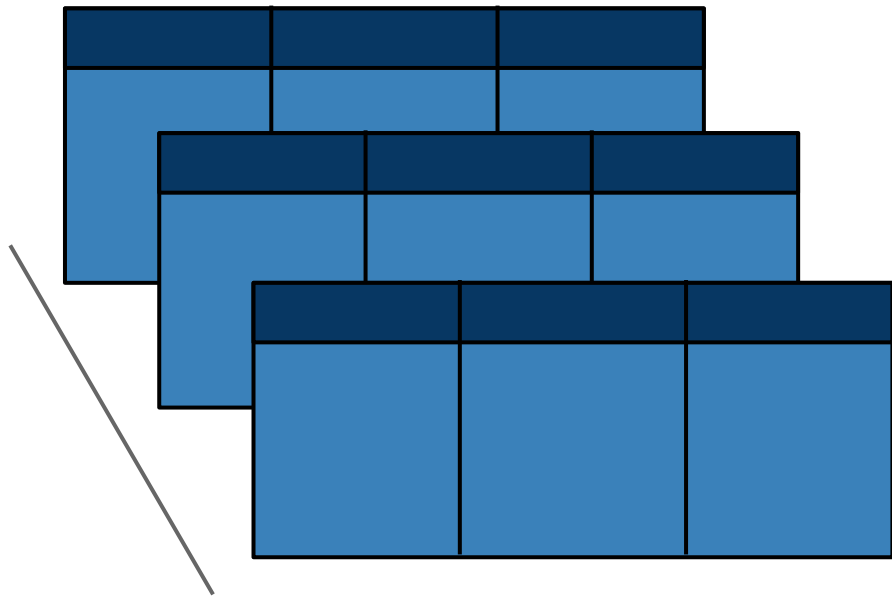
name = "TRANSACTIONS"

The diagram shows three tables stacked vertically, representing a temporary table structure. Each table has a dark blue header row and a light blue data row. The tables are offset to the right, creating a staircase effect. The top table has 3 columns, the middle table has 4 columns, and the bottom table has 3 columns.

TRANSACTIONS

The diagram shows three tables stacked vertically, representing a permanent table structure. Each table has a dark blue header row and a light blue data row. The tables are offset to the right, creating a staircase effect. The top table has 3 columns, the middle table has 4 columns, and the bottom table has 3 columns.

repartition



replace

```
to_replace = 'Dave' value = 'David'
```

| | | |
|-------|-------|-----|
| | | |
| Carol | Dave | 0.3 |
| | | |
| Bob | Carol | 0.2 |
| | | |
| from | to | amt |
| Alice | Bob | 0.1 |

| | | |
|-------|-------|-----|
| | | |
| Carol | David | 0.3 |

| | | |
|-----|-------|-----|
| | | |
| Bob | Carol | 0.2 |

| from | to | amt |
|-------|-----|-----|
| Alice | Bob | 0.1 |

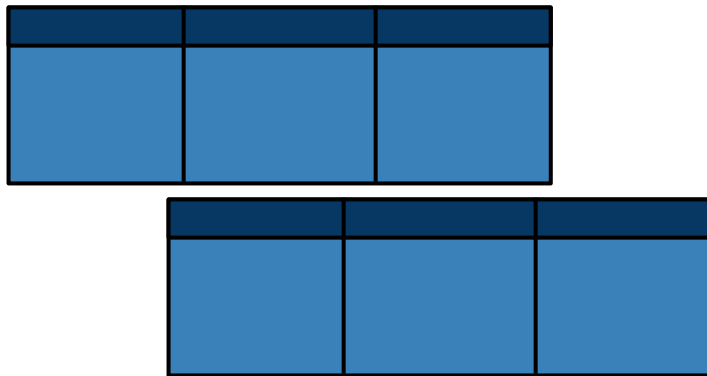
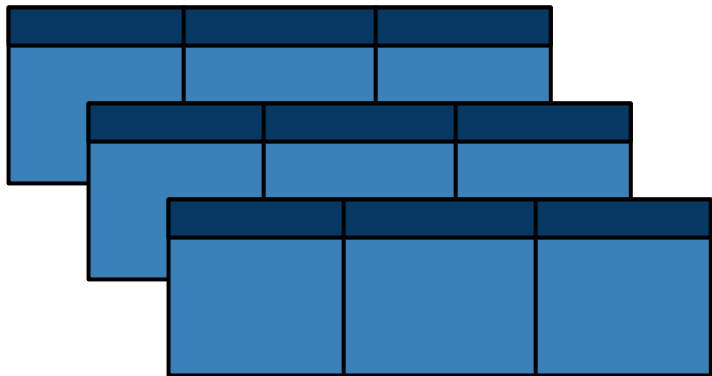
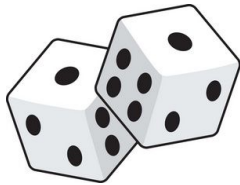
rollup

cols = ['from', 'to']

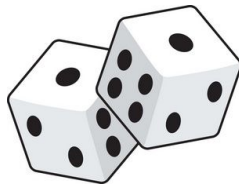
| Carol | Dave | 0.3 |
|-------|-------|-----|
| Bob | Carol | 0.2 |
| from | to | amt |
| Alice | Bob | 0.1 |

| null | null | |
|-------|-------|----------|
| Carol | null | |
| Bob | null | |
| Carol | Dave | |
| Alice | null | |
| Bob | Carol | |
| from | to | agg(amt) |
| Alice | Bob | |

sample



sampleBy



| | | |
|-------|-------|-----|
| Bob | Carol | 0.6 |
| Bob | Bob | 0.5 |
| Alice | Dave | 0.4 |
| Alice | Alice | 0.3 |
| Alice | Carol | 0.2 |
| from | to | amt |
| Alice | Bob | 0.1 |

| | | |
|-------|-------|-----|
| Bob | Carol | 0.6 |
| Bob | Bob | 0.5 |
| from | to | amt |
| Alice | Bob | 0.3 |

schema

| from | to | amt |
|------|----|-----|
| | | |

| from | to | amt |
|------|----|-----|
|------|----|-----|

select

```
cols = ['from', 'amt']
```

| from | to | amt |
|------|----|-----|
| | | |

| from | amt |
|------|-----|
| | |

selectExpr

```
expr = ["substr(from,1,1)", "amt + 10"]
```

| | | |
|-------|------|-----|
| | | |
| Carol | Dave | 0.3 |

| | | |
|-----|-------|-----|
| | | |
| Bob | Carol | 0.2 |

| from | to | amt |
|-------|-----|-----|
| Alice | Bob | 0.1 |

The diagram illustrates three overlapping tables. The top table has columns *C* and 10.3. The middle table has columns *B* and 10.2. The bottom table has columns *from* and *amt*, with values *A* and 10.1 respectively.

show

| Carol | Dave | 0.3 |
|-------|------|-----|

| Bob | Carol | 0.2 |
|-----|-------|-----|

| from | to | amt |
|-------|-----|-----|
| Alice | Bob | 0.1 |

stdout

```
+-----+-----+---+  
| from|  to|amt|  
+-----+-----+---+  
| Alice| Bob|0.1|  
| Bob|Carol|0.2|  
| Carol| Dave|0.3|  
+-----+-----+---+
```

sort

cols = ['to']

| | | | |
|-------|-------|-----|-----|
| | | | |
| Carol | Alice | 0.3 | |
| | | | |
| Bob | Carol | 0.2 | |
| | | | |
| | from | to | amt |
| | Alice | Bob | 0.1 |

| | | |
|-----|-------|-----|
| | | |
| Bob | Carol | 0.2 |

| | | |
|-------|-----|-----|
| | | |
| Alice | Bob | 0.1 |

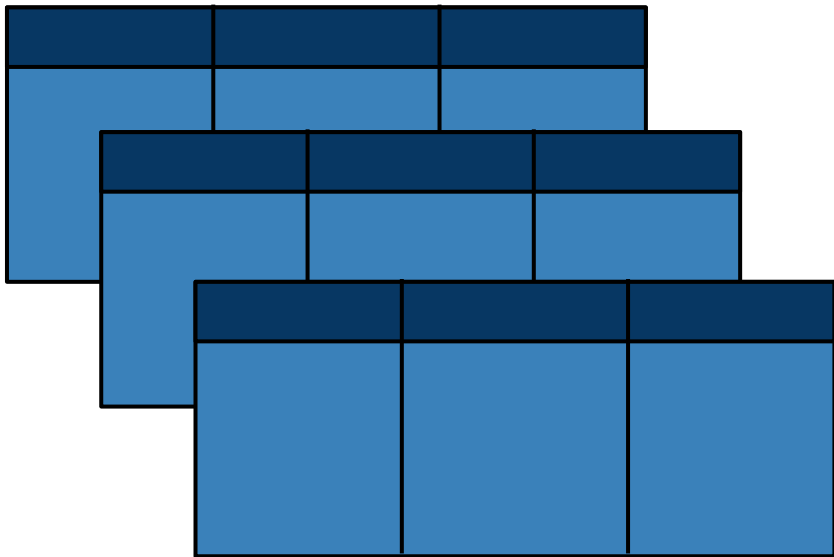
| from | to | amt |
|-------|-------|-----|
| Carol | Alice | 0.3 |

sortWithinPartitions

| Carol | Alice | 0.3 | 2 |
|-------|-------|-----|------|
| Bob | Carol | 0.2 | 2 |
| from | to | amt | p_id |
| Alice | Bob | 0.1 | 1 |

| Bob | Carol | 0.2 | 2 |
|-------|-------|-----|------|
| Carol | Alice | 0.3 | 2 |
| from | to | amt | p_id |
| Alice | Bob | 0.1 | 1 |

stat



DataFrameStatFunctions
Object
with methods: `corr`, `cov`,
`corsstab`, `freqItems`, `sampleBy`

subtract

[illegible]

| | | |
|-------|------|-----|
| | | |
| Carol | Dave | 0.1 |

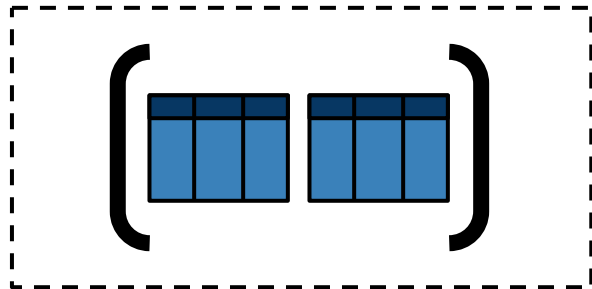
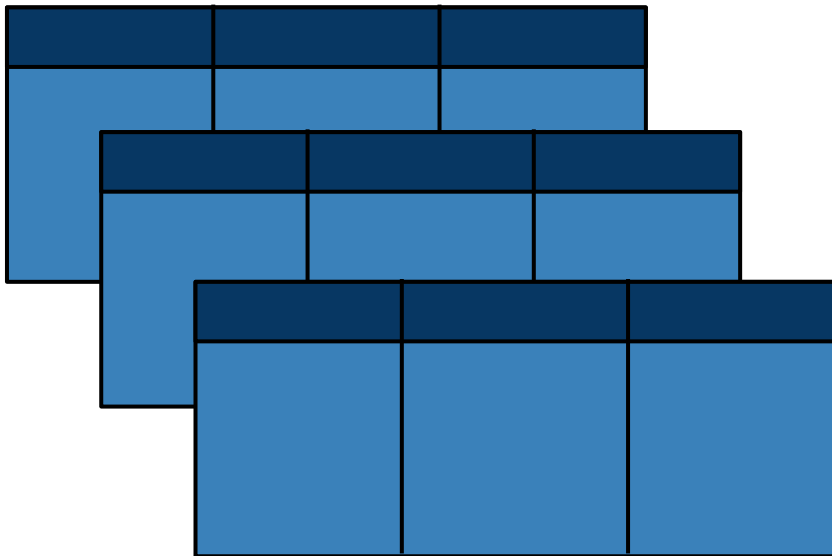
| | | |
|-----|-------|-----|
| | | |
| Bob | Carol | 0.2 |

| from | to | amt |
|-------|-----|-----|
| Alice | Bob | 0.1 |
| | Bob | 0.2 |

| from | to | amt |
|-------|------|-----|
| Carol | Dave | 0.3 |

take

num = 2



toDF

```
cols = ["seller", "buyer"]
```

| from | to | amt |
|------|----|-----|
| | | |

| seller | buyer | amt |
|--------|-------|-----|
| | | |

toJSON

| Carol | Alice | 0.3 |
|-------|-------|-----|

| Bob | Carol | 0.2 |
|-----|-------|-----|

| from | to | amt |
|-------|-----|-----|
| Alice | Bob | 0.1 |

u'{"from":"Carol","to":"Alice","amt":0.3}'

u'{"from":"Bob","to":"Carol","amt":0.2}'

u'{"from":"Alice","to":"Bob","amt":0.1}'

toPandas

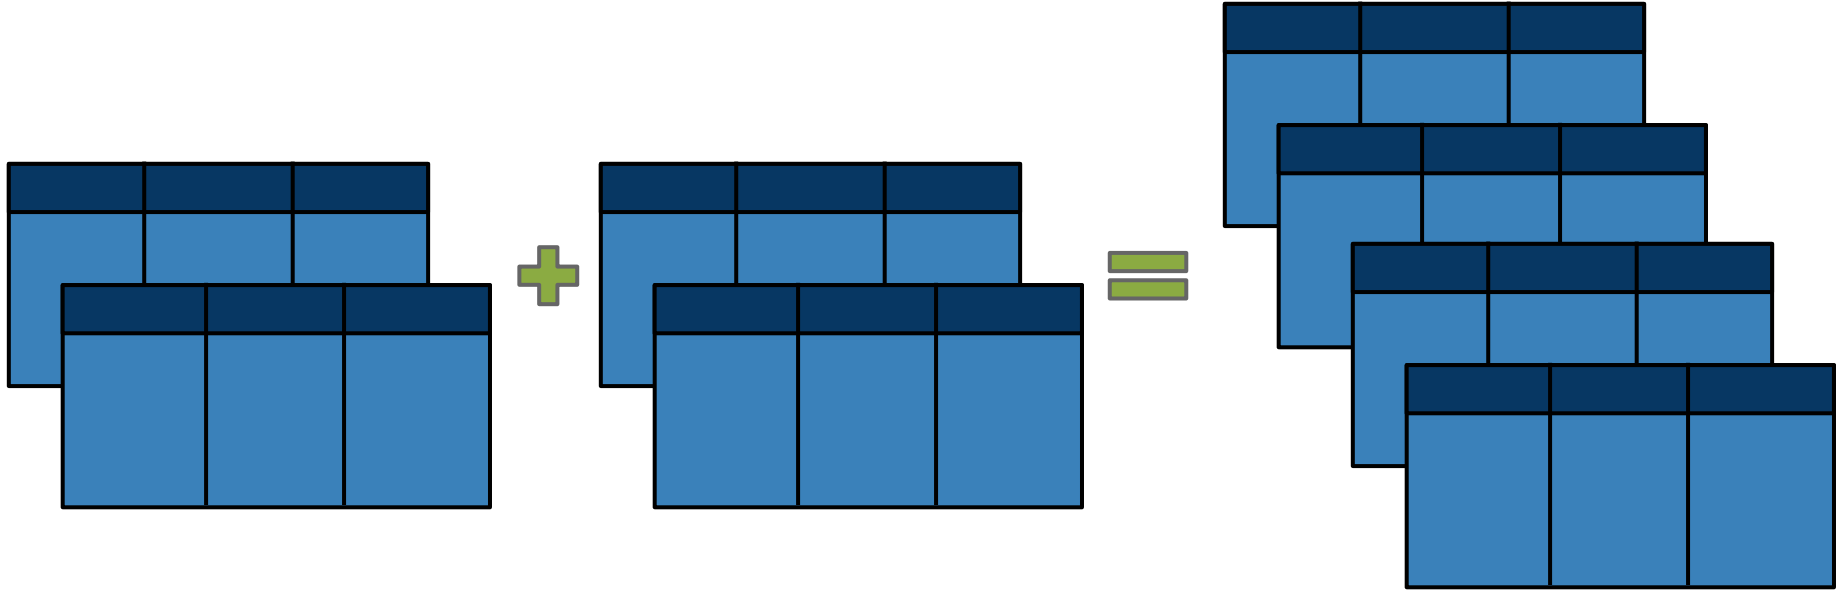
| Carol | Alice | 0.3 |
|-------|-------|-----|

| Bob | Carol | 0.2 |
|-----|-------|-----|

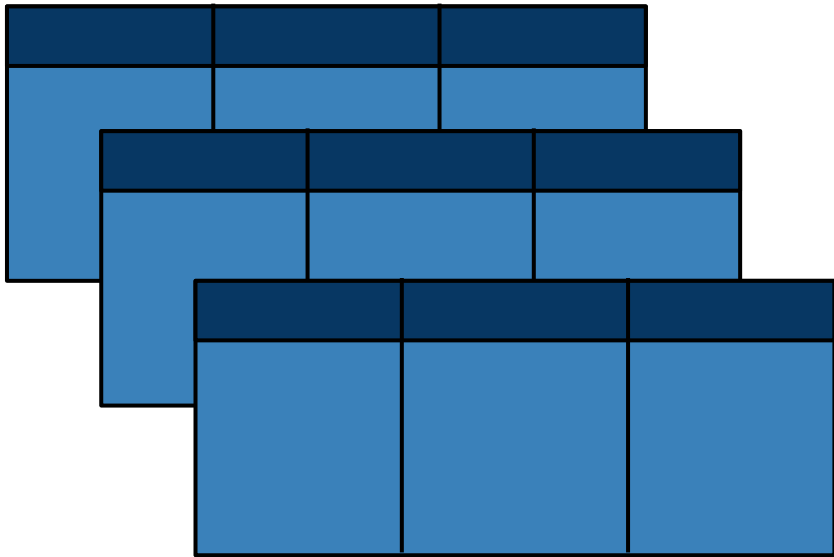
| from | to | amt |
|-------|-----|-----|
| Alice | Bob | 0.1 |

```
from to  amt
0 Alice Bob 0.1
1 Bob  Carol 0.2
2 Carol Alice 0.3
```

unionAll



unpersist



where (filter)

condition = "amt > 0.1"

| Carol | Dave | 0.3 |
|-------|-------|-----|
| | | |
| Bob | Carol | 0.2 |
| from | to | amt |
| Alice | Bob | 0.1 |

| Carol | Dave | 0.3 |
|-------|-------|-----|
| from | to | amt |
| Bob | Carol | 0.2 |

withColumn

colName = 'conf'

| from | to | amt |
|------|----|-----|
| | | |

| from | to | amt | conf |
|------|----|-----|------|
| | | | |

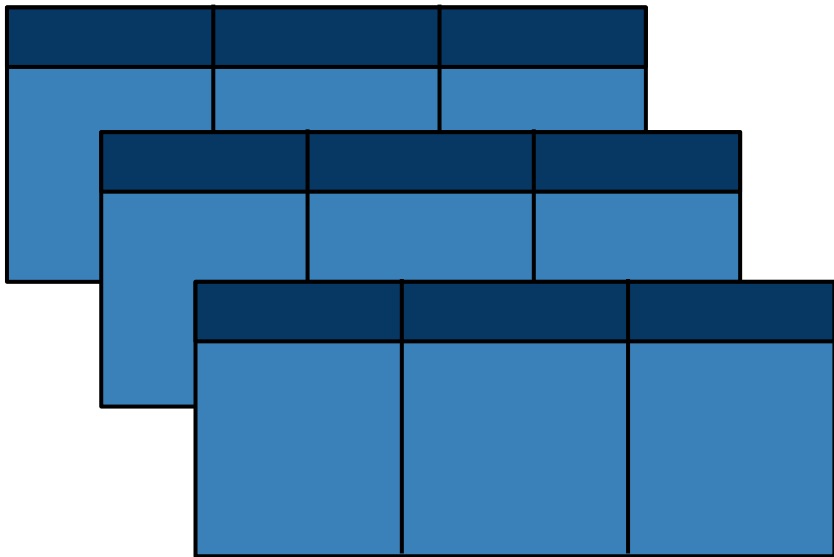
withColumnRenamed

```
existing = 'amt' col = 'amount'
```

| from | to | amt |
|------|----|-----|
| | | |

| from | to | amount |
|------|----|--------|
| | | |

write



DataFrameWriter
Object

with methods: format,
insertInto, jdbc, json, mode,
option, options, orc, parquet,
partitionBy, save, saveAsTable,
text