


# pyspark-pictures data frames


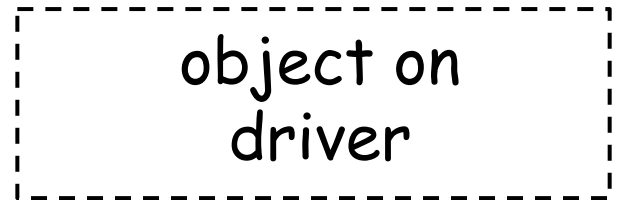


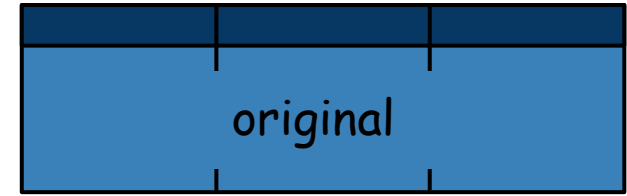
Learn the pyspark API through pictures and simple examples

<https://github.com/jkthompson/pyspark-pictures>



data frame



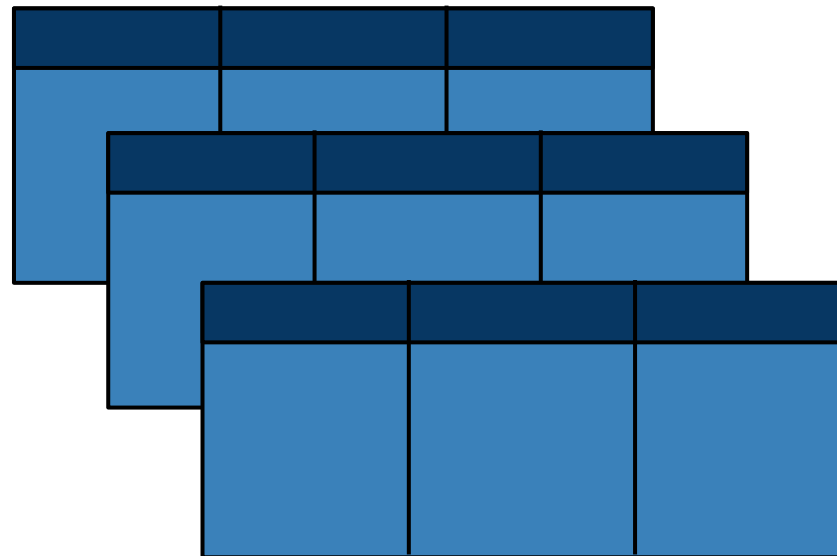
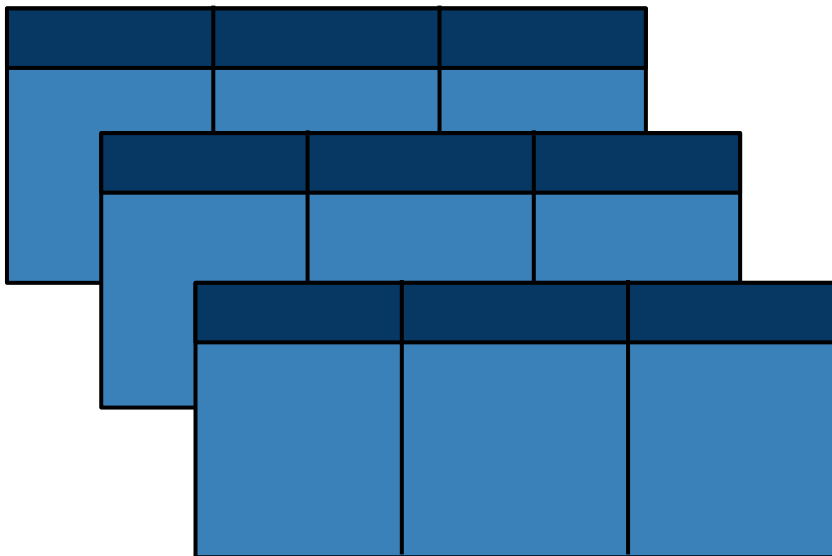
data frame row



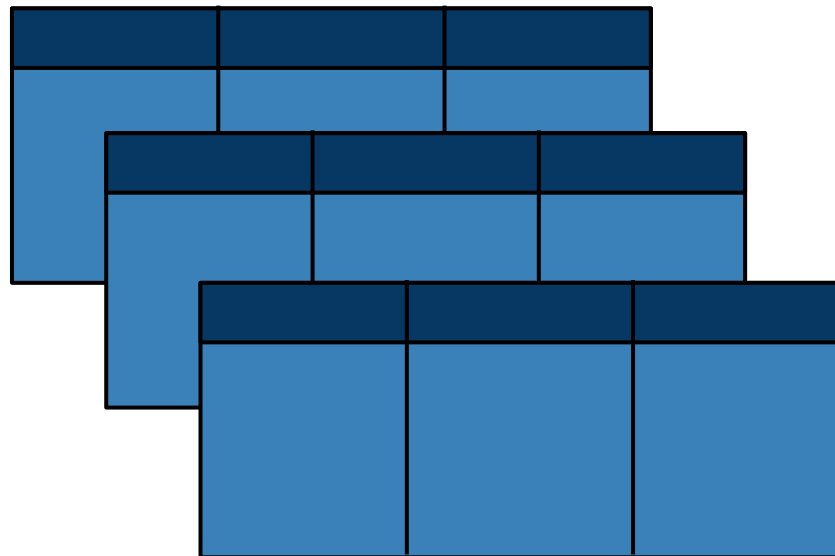
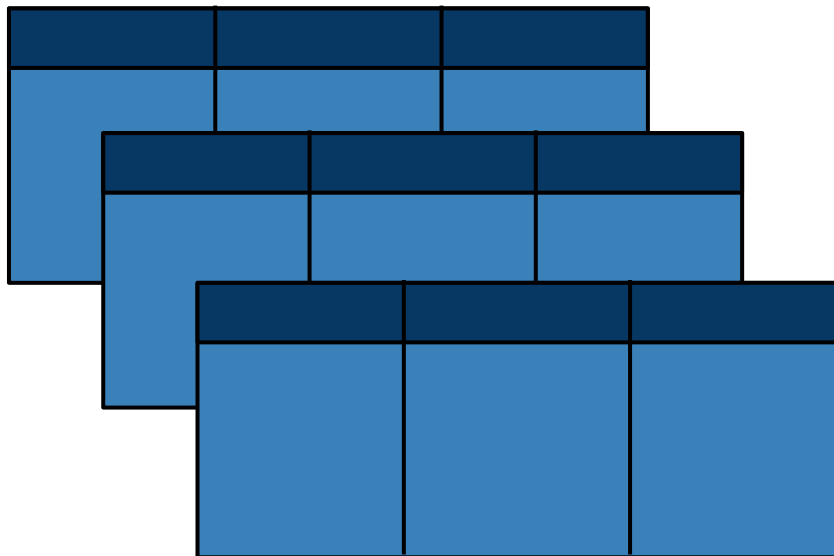
agg



# alias

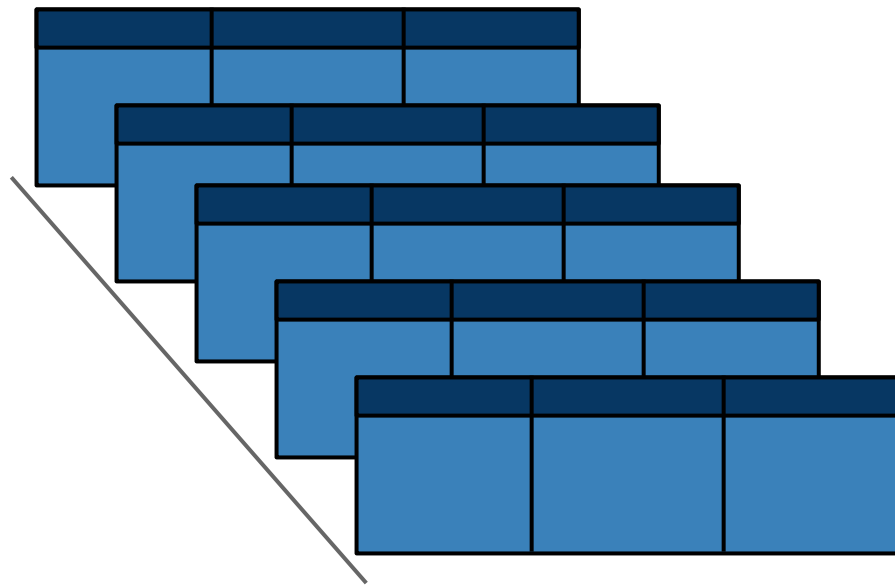
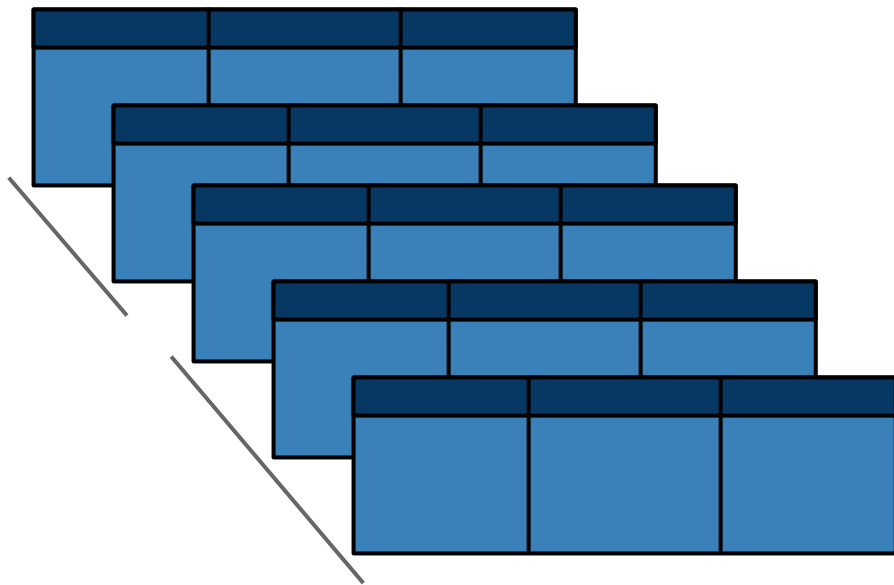


# cache


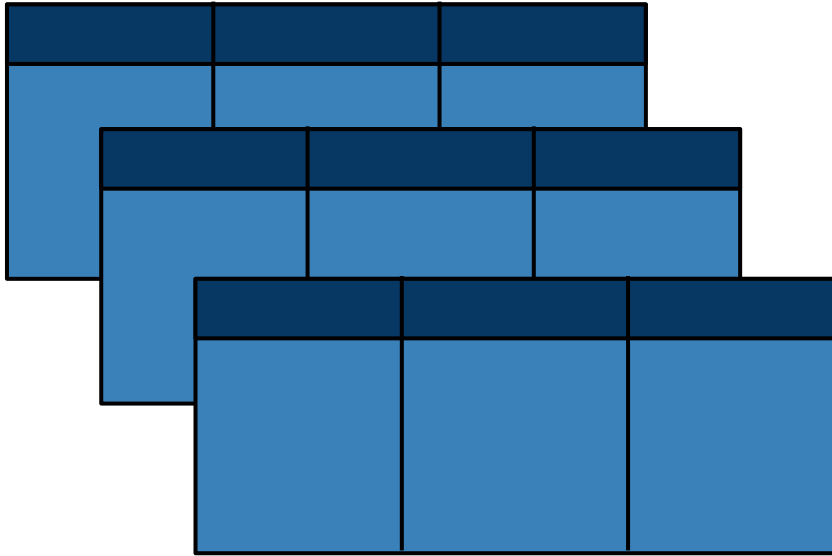


# coalesce


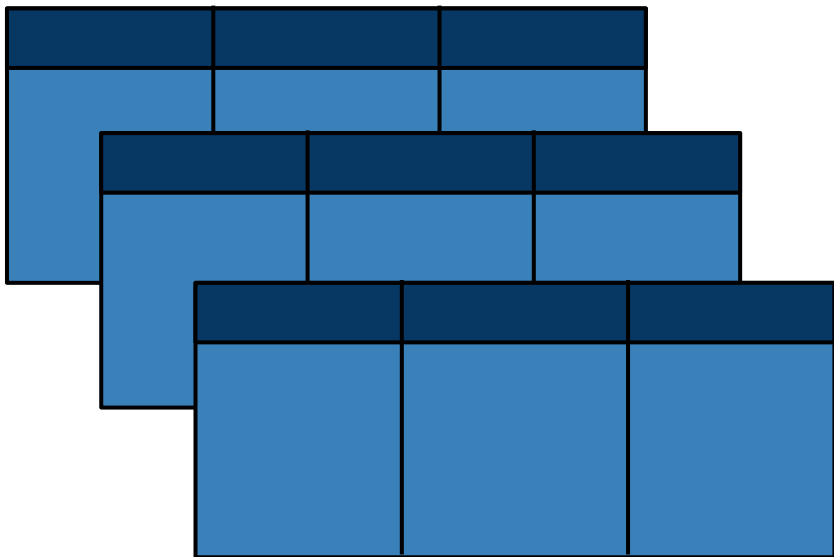
numPartitions = 1



# collect




# columns





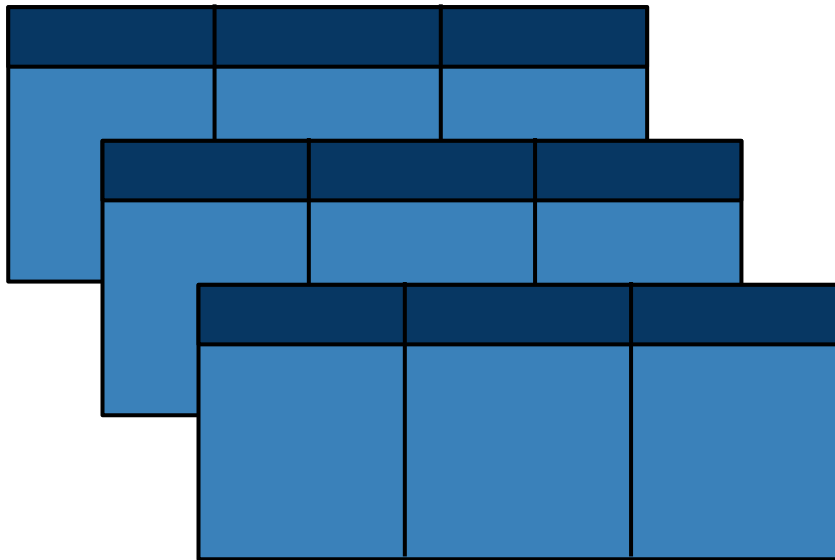
# corr



Pearson's r

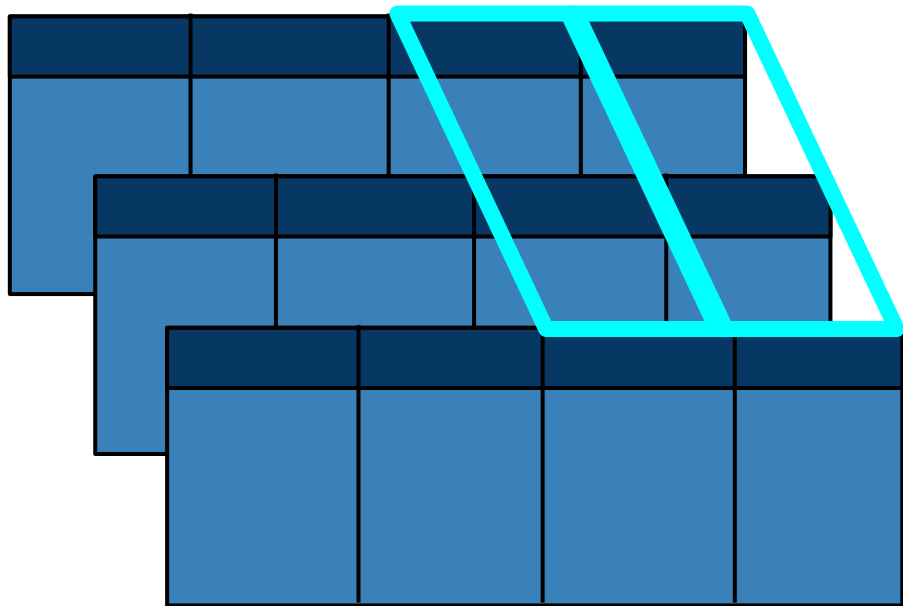
$$r = \frac{\sum_i (A_i - \bar{A})(C_i - \bar{C})}{\sqrt{\sum_i (A_i - \bar{A})^2} \sqrt{\sum_i (C_i - \bar{C})^2}}$$

# count



3

# COV

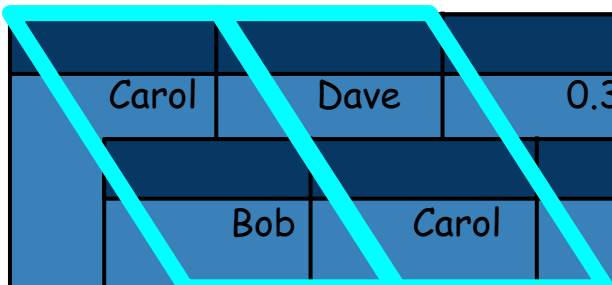


Sample Covariance

$$\frac{1}{N-1} \sum_i (A_i - \bar{A})(C_i - \bar{C})$$

# crosstab

col1 = 'from' col2 = 'to'



from	to	amt
Alice	Bob	0.1
Bob	Carol	0.2
Carol	Dave	0.3

Carol	0	0	1
Bob	0	1	0
from_to	Bob	Carol	Dave
Alice	1	0	0

# cube

\*cols = 'from', 'to'

Alice	Carol	0.2
from	to	amt
Alice	Bob	0.1

null	null	
Alice	null	
null	Carol	
null	Bob	
Alice	Carol	
from	to	agg(amt)
Alice	Bob	

# describe

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

max	
min	
stdev	
mean	
summary	amt
count	

# distinct

Bob	Carol	0.2
Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

# drop

col = 'amt'

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

Carol	Dave
Bob	Carol
from	to
Alice	Bob



# dropDuplicates

subset = ['from', 'to']

Bob	Carol	0.2
Bob	Carol	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

# dropna

how = 'any' subset = ['from', 'to']

Bob	Carol	0.2
Carol	null	0.3
Bob	Carol	Na
from	to	amt
null	Bob	0.1

Bob	Carol	0.2
from	to	amt
Bob	Carol	Na

# dtypes

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

[('from','string'), ('to','string'), ('amt','double')]

# explain

extended = True

Carol	Dave	0.3

Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

== Parsed Logical Plan ==

...

== Analyzed Logical Plan ==

...

== Optimized Logical Plan ==

...

== Physical Plan ==

...

== RDD ==

# fillna

value = 'unknown' subset = ['from', 'to']

Carol	null	0.3

Bob	Carol	nan

from	to	amt
null	Bob	0.1

Carol	unknown	0.3

Bob	Carol	nan

from	to	amt
unknown	Bob	0.1

# filter

condition = "amt > 0.1"

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1


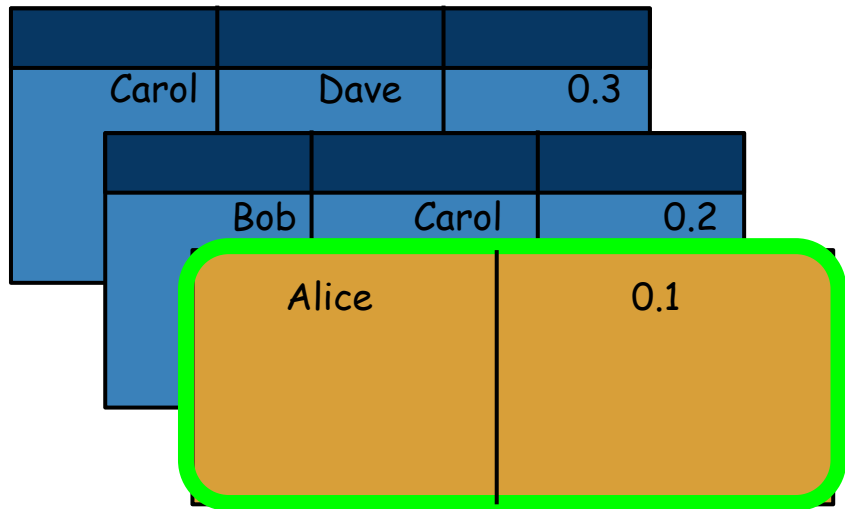
Carol	Dave	0.3
from	to	amt
Bob	Carol	0.2

# first

	Carol	Dave 0.3
	Bob	Carol 0.2
	from	to amt
	Alice	Bob 0.1

Row(from='Alice', to='Bob', amt=0.1)

# flatMap





# foreach

Carol	Dave	0.3

Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

side effects  
(e.g print) ←

\*no return value,  
original DataFrame  
unchanged

# foreachPartition

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

side effects  
(e.g print) ←

\*no return value,  
original DataFrame  
unchanged


# freqItems

cols = ['from', 'amt'] support = 0.8

Carol	Bob	0.1
Alice	Bob	0.5
Alice	Bob	0.1
Alice	Dave	0.1
from	to	amt
Bob	Carol	0.1

from_freqItems	amt_freqItems
[Alice]	[0.1]

# groupBy (groupby)



*GroupedData Object*  
with methods: *agg, avg, count,*  
*max, mean, min, pivot, sum*

# groupBy(col1).avg(col2)

col1 = 'from' col2 = 'amt'

Carol	Dave	0.3
Alice	Carol	0.2
from	to	amt
Alice	Bob	0.1

Carol	0.3
from	avg(amt)
Alice	0.15

# head

n = 2

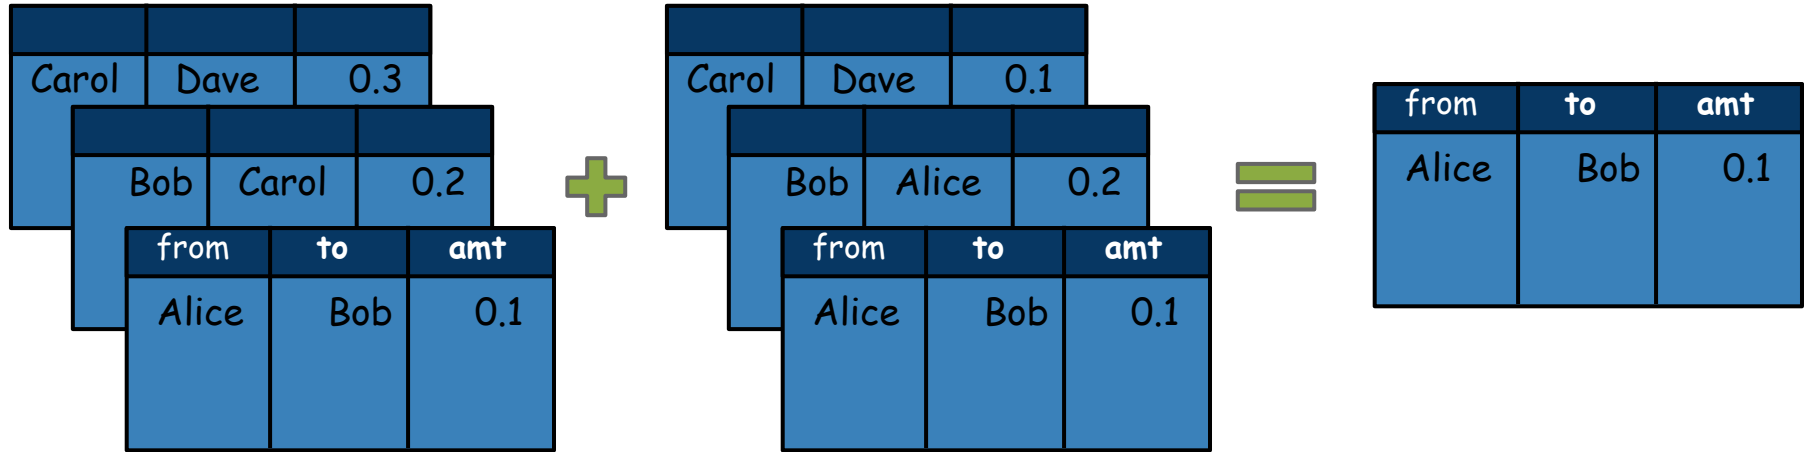
Carol	Dave	0.3

Bob	Carol	0.2


from	to	amt
Alice	Bob	0.1

[Row(from=u'Alice', to=u'Bob', amt=0.1),  
Row(from=u'Bob', to=u'Carol', amt=0.2)]

# intersect



# isLocal



False




# join

joinExprs = x.to==y.name  
joinType = 'inner'

Carol	Dave	0.3

Bob	Carol	0.2


from	to	amt
Alice	Bob	0.1



Dave	80

Bob	40

name	age
Alice	20





Alice	Bob	0.1	Bob	40


from	to	amt	name	age
Carol	Dave	0.3	Dave	80

# limit

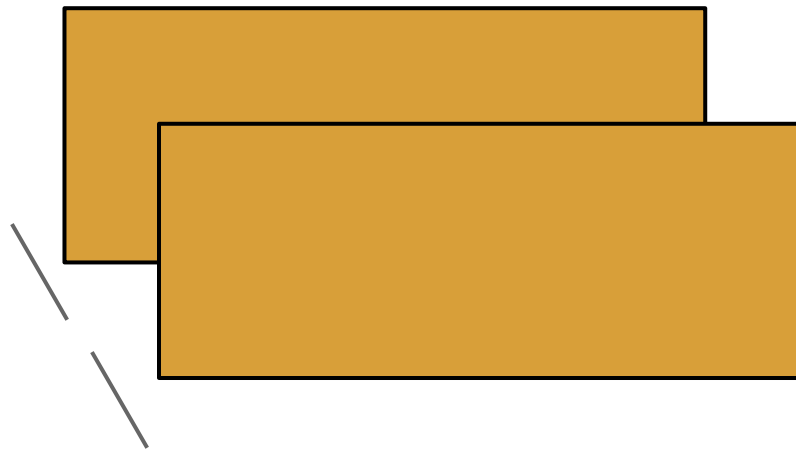

num = 2



# map



# mapPartitions



# na

Bob	Carol	0.2
Carol	null	0.3
Bob	Carol	Na
from	to	amt
null	Bob	0.1

## DataFrameNaFunction

s

Bob	Carol	0.2
Carol	null	0.3
Bob	Carol	Na
from	to	amt
null	Bob	0.1

# orderBy

```
cols = ['from'], ascending = [False]
```

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

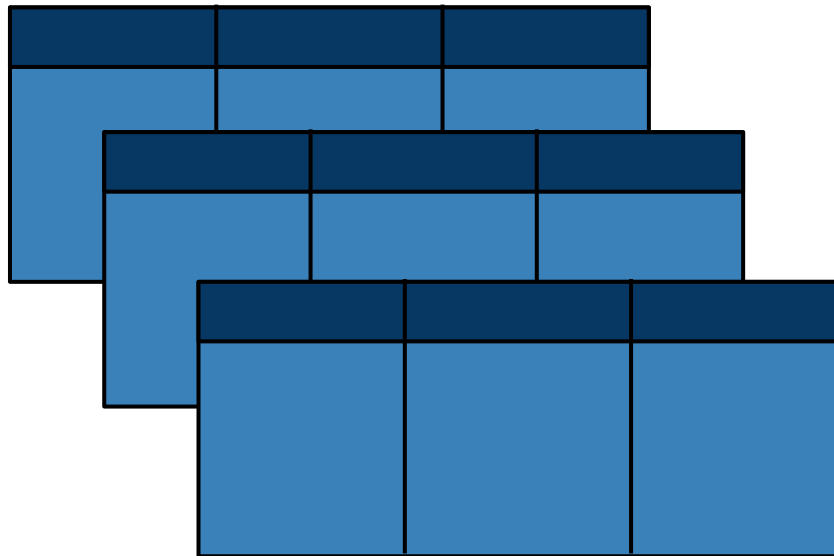
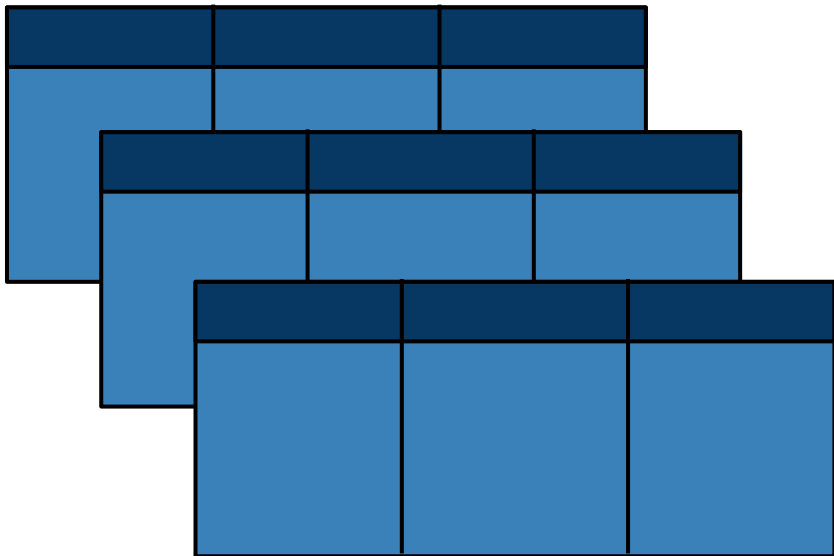
Alice	Bob	0.1

Bob	Carol	0.2

from	to	amt
Carol	Dave	0.3

# persist

```
storageLevel =  
StorageLevel(MEMORY_ONLY_SER)
```



# printSchema

Carol	Dave	0.3

Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

stdout

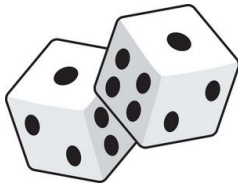
root

|-- from: string (nullable = true)  
|-- to: string (nullable = true)  
|-- amt: double (nullable = true)



# randomSplit

weights = [0.5,0.5]

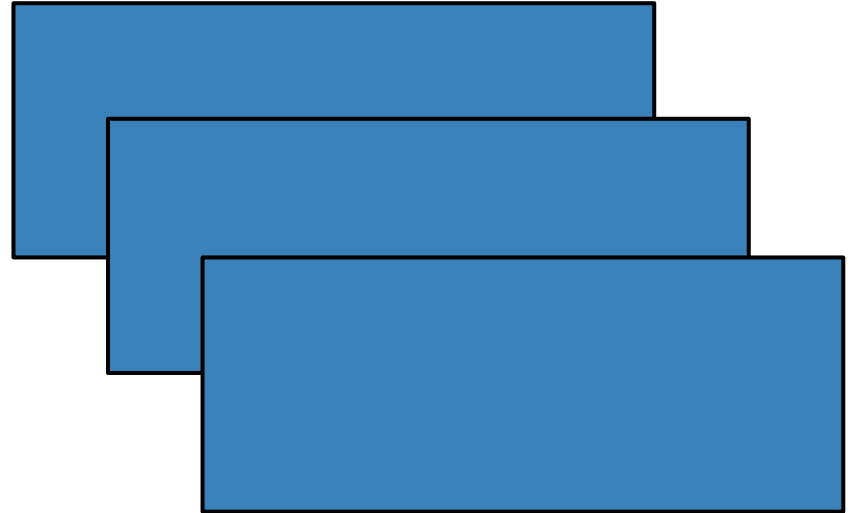
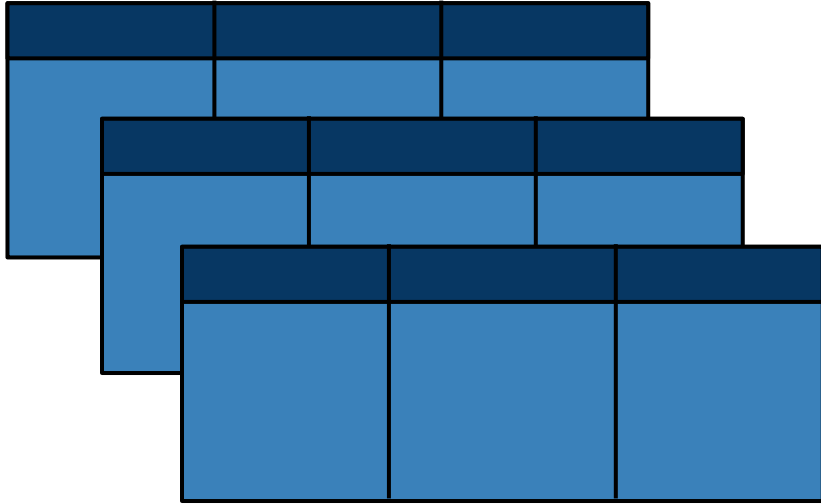


Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

Carol	Dave	0.3
from	to	amt
Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

# rdd



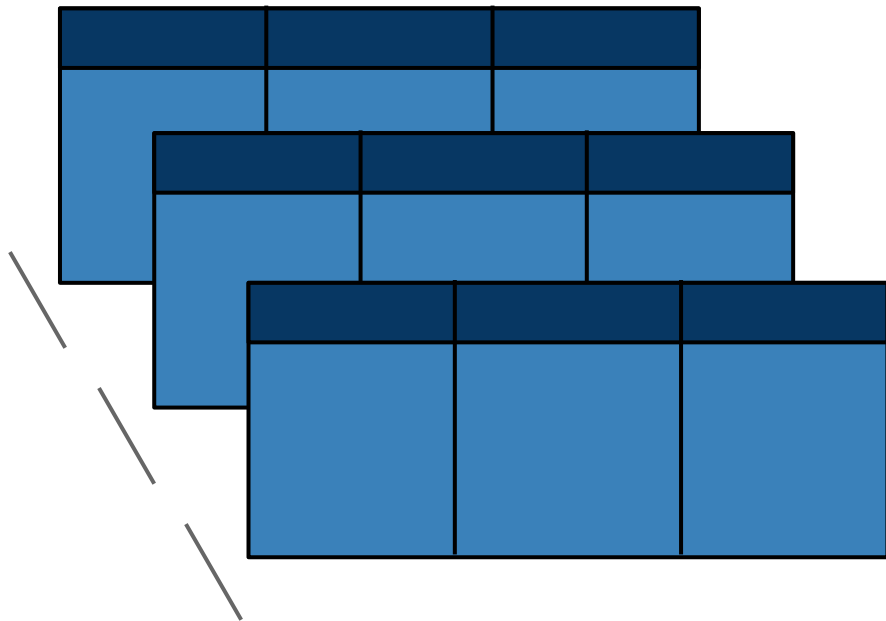
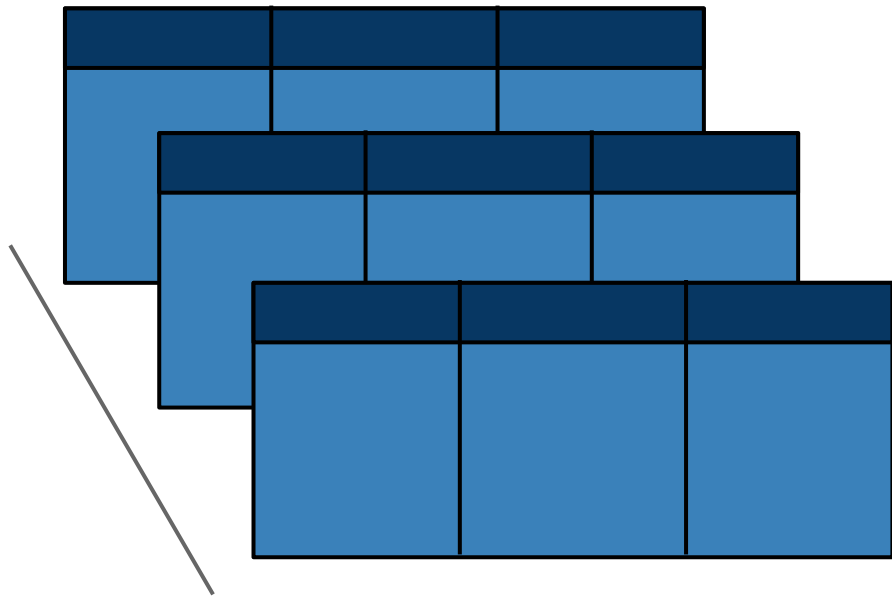
# registerTempTable

name = "TRANSACTIONS"



TRANSACTIONS

# repartition



# replace

```
to_replace = 'Dave' value = 'David'
```

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

Carol	David	0.3

Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

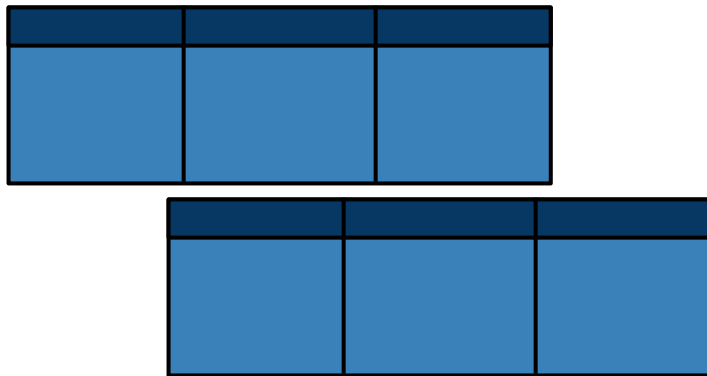
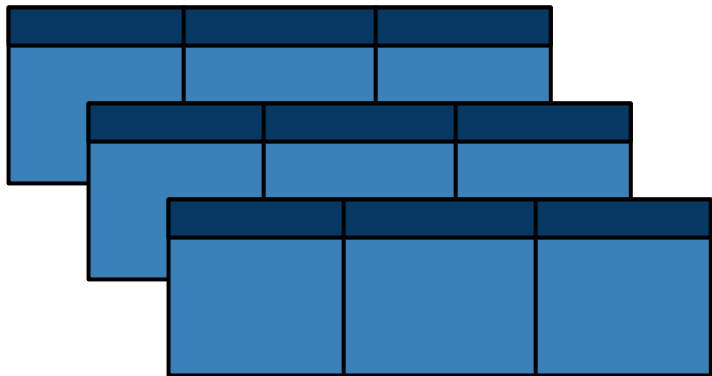
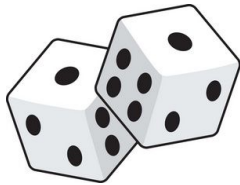
# rollup

cols = ['from', 'to']

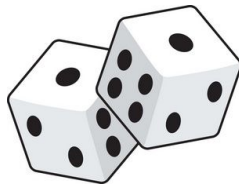
Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

null	null	
Carol	null	
Bob	null	
Carol	Dave	
Alice	null	
Bob	Carol	
from	to	agg(amt)
Alice	Bob	

# sample



# sampleBy



Bob	Carol	0.6
Bob	Bob	0.5
Alice	Dave	0.4
Alice	Alice	0.3
Alice	Carol	0.2
from	to	amt
Alice	Bob	0.1

Bob	Carol	0.6
Bob	Bob	0.5
from	to	amt
Alice	Bob	0.3



# schema

from	to	amt

from	to	amt
------	----	-----

# select

```
cols = ['from', 'amt']
```

from	to	amt

from	amt

# selectExpr

```
expr = ["substr(from,1,1)", "amt + 10"]
```

Carol	Dave	0.3

Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

The diagram illustrates three overlapping tables, each with two columns: 'from' and 'amt'. The tables are arranged in a descending staircase pattern, with each subsequent table shifted further to the right and down. The top table has 'A' in the 'from' column and '10.1' in the 'amt' column. The middle table has 'B' in the 'from' column and '10.2' in the 'amt' column. The bottom table has 'C' in the 'from' column and '10.3' in the 'amt' column. The overlapping areas of the tables are shaded in a darker blue, while the non-overlapping areas are a lighter blue.

from	amt
A	10.1

from	amt
B	10.2

from	amt
C	10.3

# show

Carol	Dave	0.3

Bob	Carol	0.2

from	to	amt
Alice	Bob	0.1

stdout

```
+-----+-----+-----+
| from|  to|amt|
+-----+-----+-----+
| Alice| Bob|0.1|
|  Bob| Carol|0.2|
| Carol| Dave|0.3|
+-----+-----+-----+
```

# sort

cols = ['to']

Carol	Alice	0.3	
Bob	Carol	0.2	
	from	to	amt
	Alice	Bob	0.1

Bob	Carol	0.2

Alice	Bob	0.1


from	to	amt
Carol	Alice	0.3

# sortWithinPartitions

Carol	Alice	0.3	2
Bob	Carol	0.2	2
from	to	amt	p_id
Alice	Bob	0.1	1

Bob	Carol	0.2	2
Carol	Alice	0.3	2
from	to	amt	p_id
Alice	Bob	0.1	1

# stat



# stat



# subtract

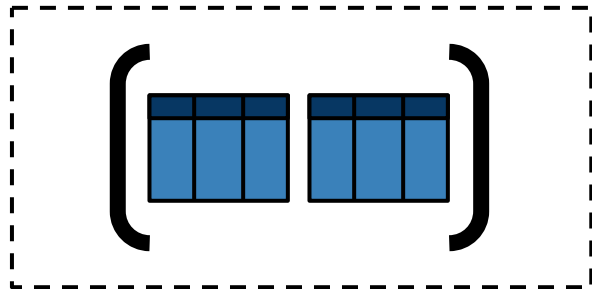
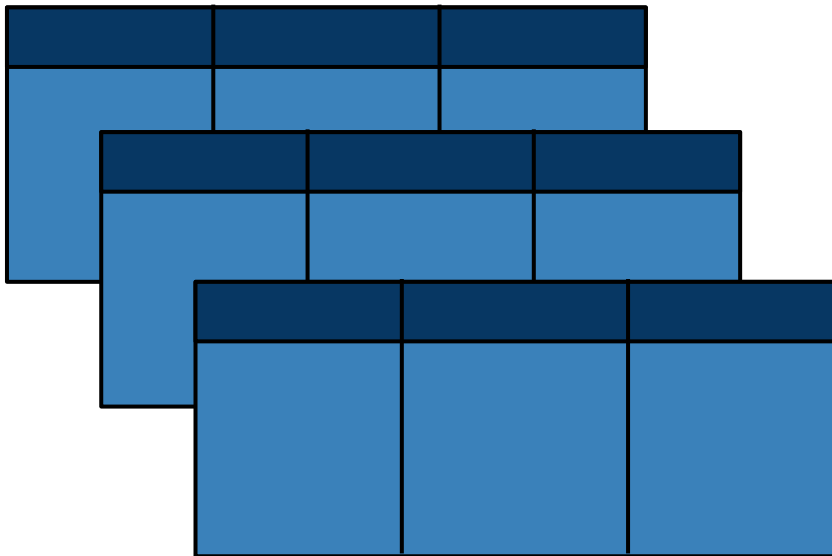
from	to	amt
Carol	Dave	0.3
Bob	Carol	0.2
Alice	Bob	0.1

from	to	amt
Carol	Dave	0.1
Bob	Carol	0.2
Alice	Bob	0.1

from	to	amt
Carol	Dave	0.3

# take

num = 2



# toDF

```
cols = ["seller", "buyer"]
```

from	to	amt

seller	buyer	amt

# toJSON

cols = ['to']

Carol	Alice	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

u'{"from":"Carol","to":"Alice","amt":0.3}'

u'{"from":"Bob","to":"Carol","amt":0.2}'

u'{"from":"Alice","to":"Bob","amt":0.1}'

# toPandas


Carol	Alice	0.3

Bob	Carol	0.2

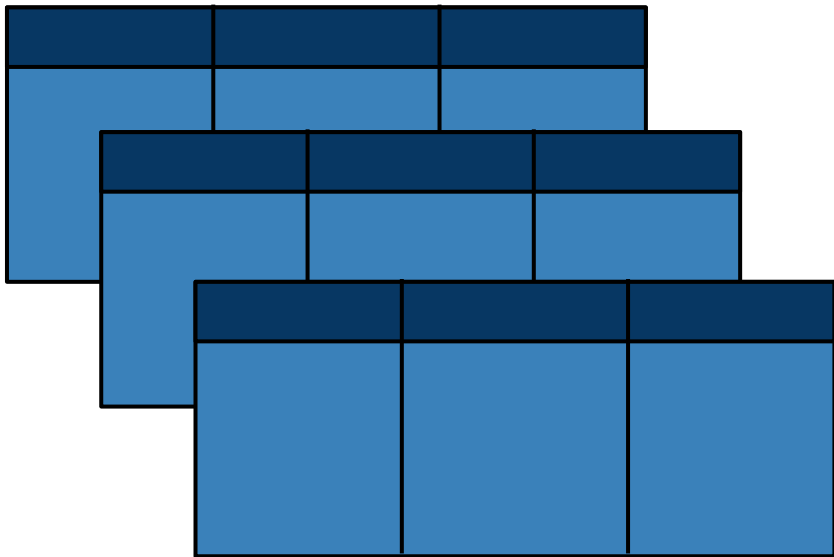
from	to	amt
Alice	Bob	0.1

```
from to  amt
0 Alice Bob 0.1
1 Bob  Carol 0.2
2 Carol Alice 0.3
```

# unionAll



# unpersist



# where (filter)

condition = "amt > 0.1"

Carol	Dave	0.3
Bob	Carol	0.2
from	to	amt
Alice	Bob	0.1

Carol	Dave	0.3
from	to	amt
Bob	Carol	0.2



# withColumn

```
colName = 'conf'
```

from	to	amt

from	to	amt	conf

# withColumnRenamed

```
existing = 'amt' col = 'amount'
```

from	to	amt

from	to	amount

# write

